

# TailRisk

Risk-Aware Machine Learning  
for Tail Risk Modeling

Technical Documentation

Version 0.1.2

Vishal Lakshmi Narayanan  
[lvishal1607@gmail.com](mailto:lvishal1607@gmail.com)  
[GitHub Repository](#)

November 22, 2025

## Abstract

TailRisk is a Python package designed for building machine learning models that excel at predicting extreme outcomes in insurance claims, financial losses, and other tail-risk scenarios. Traditional machine learning models optimize for average performance (MSE, MAE), which fails catastrophically when predicting rare, extreme events. This package implements novel methodologies including Loss-at-Risk (LaR) regression, CVaR-weighted ensembles, and hybrid meta-learning to address these critical challenges in risk modeling.

**Keywords:** Tail Risk, Machine Learning, Insurance, Finance, CVaR, Extreme Value Prediction, Quantile Regression

# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	The Tail Risk Problem . . . . .	5
1.2	Why MSE Fails . . . . .	5
1.3	The TailRisk Solution . . . . .	5
<b>2</b>	<b>Installation</b>	<b>6</b>
2.1	Requirements . . . . .	6
2.2	Installation via PyPI . . . . .	6
2.3	Installation from Source . . . . .	6
2.4	Verification . . . . .	6
<b>3</b>	<b>Quick Start</b>	<b>7</b>
3.1	Basic Example: LaR Regressor . . . . .	7
3.2	Advanced Example: Hybrid Meta-Learner . . . . .	7
<b>4</b>	<b>API Reference</b>	<b>9</b>
4.1	Models . . . . .	9
4.1.1	LaRRegressor . . . . .	9
4.1.2	HybridMetaLearner . . . . .	10
4.1.3	CVaRWeightedEnsemble . . . . .	11
4.2	Metrics . . . . .	12
4.2.1	Conditional Value-at-Risk (CVaR) . . . . .	12
4.2.2	Tail Coverage Ratio (TCR) . . . . .	12
4.2.3	Detection Rate . . . . .	13
4.2.4	tail_validation_summary . . . . .	13
<b>5</b>	<b>Methodology</b>	<b>14</b>
5.1	The Hybrid Meta-Learning Framework . . . . .	14
5.1.1	Stage 1: Diverse Base Models . . . . .	14
5.1.2	Stage 2a: Quantile Meta-Model . . . . .	14
5.1.3	Stage 2b: LaR-Weighted Optimization . . . . .	14
5.1.4	Stage 3: Hybrid Blending . . . . .	14
<b>6</b>	<b>Use Cases</b>	<b>15</b>
6.1	Insurance Claims Prediction . . . . .	15
6.2	Financial Risk Management . . . . .	15
6.3	Healthcare Cost Prediction . . . . .	16
<b>7</b>	<b>Model Comparison Example</b>	<b>17</b>
7.1	Complete Workflow . . . . .	17
7.2	Expected Output . . . . .	18
<b>8</b>	<b>Best Practices</b>	<b>19</b>
8.1	Model Selection . . . . .	19
8.2	Hyperparameter Tuning . . . . .	19
8.2.1	LaRRegressor Alpha Parameter . . . . .	19
8.2.2	HybridMetaLearner Configuration . . . . .	19
8.3	Data Considerations . . . . .	20
8.3.1	Minimum Sample Size . . . . .	20

8.3.2 Train/Test Split Strategy . . . . .	20
8.4 Evaluation . . . . .	20
<b>9 Scikit-learn Integration</b>	<b>21</b>
9.1 GridSearchCV . . . . .	21
9.2 Pipeline . . . . .	21
9.3 Cross-Validation . . . . .	21
<b>10 Advanced Topics</b>	<b>23</b>
10.1 Custom Base Estimators . . . . .	23
10.2 Ensemble Stacking . . . . .	23
10.3 Production Deployment . . . . .	23
10.3.1 Model Serialization . . . . .	23
10.3.2 Monitoring . . . . .	24
<b>11 Troubleshooting</b>	<b>25</b>
11.1 Common Issues . . . . .	25
11.1.1 Issue: Negative Predictions . . . . .	25
11.1.2 Issue: Poor Tail Performance . . . . .	25
11.1.3 Issue: Training Too Slow . . . . .	25
11.2 Debugging . . . . .	26
<b>12 Theoretical Background</b>	<b>27</b>
12.1 Quantile Regression . . . . .	27
12.2 Extreme Value Theory . . . . .	27
12.3 CVaR as Coherent Risk Measure . . . . .	27
<b>13 References</b>	<b>28</b>
13.1 Package Documentation . . . . .	28
13.2 Related Software . . . . .	28
<b>14 Citation</b>	<b>29</b>
<b>15 License</b>	<b>29</b>
<b>16 Contact &amp; Support</b>	<b>30</b>
16.1 Author . . . . .	30
16.2 Bug Reports . . . . .	30
16.3 Feature Requests . . . . .	30
16.4 Contributing . . . . .	30

## 1 Introduction

### 1.1 The Tail Risk Problem

Traditional machine learning models optimize metrics such as Mean Squared Error (MSE) or Mean Absolute Error (MAE). While these metrics work well for average performance, they treat all prediction errors equally. This approach is **fundamentally flawed** for tail risk applications:

- **Insurance claims:** Predicting \$500 for a claim that costs \$50,000 creates severe financial exposure
- **Financial risk:** Underestimating tail risk leads to inadequate capital reserves (2008 financial crisis)
- **Healthcare costs:** Missing catastrophic cases can bankrupt risk pools

### 1.2 Why MSE Fails

Consider the mathematical formulation of MSE:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (1)$$

In this formulation, a \$10,000 error on a \$500 claim receives the *same weight* as a \$10,000 error on a \$100,000 claim. For tail risk applications, this is catastrophic.

### 1.3 The TailRisk Solution

TailRisk implements three key innovations:

1. **Loss-at-Risk (LaR) Regression:** Weighted regression with adaptive importance based on claim magnitude
2. **Hybrid Meta-Learning:** Two-stage architecture combining quantile regression and LaR optimization
3. **Tail-Focused Metrics:** CVaR, Tail Coverage Ratio, and Detection Rate for proper evaluation

## 2 Installation

### 2.1 Requirements

- Python 3.8 or higher
- NumPy  $\geq$  1.21.0
- Pandas  $\geq$  1.3.0
- Scikit-learn  $\geq$  1.0.0
- SciPy  $\geq$  1.7.0
- Matplotlib  $\geq$  3.4.0

### 2.2 Installation via PyPI

The recommended installation method is via pip:

```
1 pip install tailrisk
```

### 2.3 Installation from Source

For development or latest features:

```
1 git clone https://github.com/VishalLakshmiNarayanan/TailriskLib.git
2 cd TailriskLib
3 pip install -e .
```

### 2.4 Verification

Verify successful installation:

```
1 import tailrisk
2 print(f"TailRisk version: {tailrisk.__version__}")
```

## 3 Quick Start

### 3.1 Basic Example: LaR Regressor

```

1 import numpy as np
2 from sklearn.model_selection import train_test_split
3 from tailrisk import LaRRegressor, tail_validation_summary
4
5 # Generate heavy-tailed data
6 X = np.random.randn(1000, 10)
7 y = np.random.exponential(scale=1000, size=1000)
8
9 # Split data
10 X_train, X_test, y_train, y_test = train_test_split(
11     X, y, test_size=0.2, random_state=42
12 )
13
14 # Train LaR model
15 model = LaRRegressor(alpha=2.0)
16 model.fit(X_train, y_train)
17
18 # Predict and evaluate
19 predictions = model.predict(X_test)
20 metrics = tail_validation_summary(y_test, predictions)
21
22 print(f"CVaR(95%): ${metrics['cvar_95']:.2f}")
23 print(f"Tail Coverage Ratio: {metrics['tcr_99']:.3f}")

```

### 3.2 Advanced Example: Hybrid Meta-Learner

```

1 from tailrisk import HybridMetaLearner
2 from sklearn.ensemble import (
3     RandomForestRegressor,
4     GradientBoostingRegressor
5 )
6
7 # Define diverse base models
8 base_models = [
9     ('rf', RandomForestRegressor(
10         n_estimators=100,
11         max_depth=10,
12         random_state=42
13     )),
14     ('gb', GradientBoostingRegressor(
15         n_estimators=100,
16         max_depth=5,
17         random_state=42
18     ))
19 ]
20
21 # Create hybrid meta-learner
22 model = HybridMetaLearner(
23     base_estimators=base_models,
24     quantile=0.95,           # Focus on 95th percentile
25     blend_lambda=0.25,       # 25% quantile, 75% LaR
26     lar_alpha=1.5,          # LaR weighting strength
27     cv_folds=5              # Cross-validation folds
28 )

```

```
29  
30 # Train and predict  
31 model.fit(X_train, y_train)  
32 predictions = model.predict(X_test)
```

## 4 API Reference

### 4.1 Models

#### 4.1.1 LaRRegressor

**Class:** tailrisk.LaRRegressor

Loss-at-Risk weighted regression model.

**Parameters:**

- **alpha** : *float*, default=2.0
  - Weight scaling factor
  - Higher values increase focus on large claims
  - Recommended range: 1.0 to 3.0
- **base\_estimator** : *estimator*, default=None
  - Base regression model
  - Must support `sample_weight` parameter
  - If None, uses `LinearRegression()`

**Methods:**

- `fit(X, y, sample_weight=None)` : Fit the model
- `predict(X)` : Generate predictions

**Attributes:**

- `base_estimator_` : Fitted base estimator

#### Mathematical Formulation:

The LaR objective function is defined as:

$$\mathcal{L}_{\text{LaR}} = \frac{1}{n} \sum_{i=1}^n w_i (y_i - \hat{y}_i)^2 \quad (2)$$

where the weights are:

$$w_i = 1 + \alpha \cdot \frac{y_i}{\max(y)} \quad (3)$$

This ensures larger target values receive proportionally higher importance during optimization.

#### 4.1.2 HybridMetaLearner

**Class:** `tailrisk.HybridMetaLearner`

Advanced two-stage meta-learning ensemble for tail risk prediction.

**Architecture:**

1. **Stage 1:** Generate meta-features using cross-validated predictions from diverse base models
2. **Stage 2a:** Train quantile regression meta-model focusing on high quantiles
3. **Stage 2b:** Optimize LaR-weighted combination of base model predictions
4. **Stage 3:** Blend quantile and LaR predictions for final output

**Parameters:**

- **base\_estimators** : *list of (str, estimator) tuples*
  - List of base models to ensemble
  - Recommended: Use diverse model types (trees, linear, boosting)
- **quantile** : *float*, default=0.95
  - Target quantile for quantile regression
  - Range: 0.90 to 0.99
- **blend\_lambda** : *float*, default=0.25
  - Blending weight:  $\hat{y} = \lambda \cdot \hat{y}_{\text{quantile}} + (1 - \lambda) \cdot \hat{y}_{\text{LaR}}$
  - Range: 0.0 to 1.0
  - Recommended: 0.2 to 0.3
- **lar\_alpha** : *float*, default=1.5
  - LaR weight scaling factor
  - Range: 1.0 to 2.0
- **cv\_folds** : *int*, default=5
  - Number of cross-validation folds
  - Typical range: 3 to 10

**Final Prediction Formula:**

$$\hat{y}_{\text{final}} = \lambda \cdot \hat{y}_{\text{quantile}} + (1 - \lambda) \cdot \hat{y}_{\text{LaR}} \quad (4)$$

#### 4.1.3 CVaRWeightedEnsemble

**Class:** tailrisk.CVaRWeightedEnsemble

Ensemble that weights models based on inverse CVaR performance.

**Parameters:**

- **estimators** : *list of (str, estimator) tuples*
- **alpha** : *float*, default=0.95
  - CVaR percentile for weight calculation
  - Models with lower CVaR receive higher weight

## 4.2 Metrics

### 4.2.1 Conditional Value-at-Risk (CVaR)

**Function:** `tailrisk.cvar_loss(y_true, y_pred, alpha=0.95)`

**Definition:**

CVaR measures the average squared error in the worst  $(1 - \alpha)\%$  of predictions:

$$\text{CVaR}_\alpha = \mathbb{E}[(Y - \hat{Y})^2 \mid (Y - \hat{Y})^2 \geq \text{VaR}_\alpha] \quad (5)$$

where  $\text{VaR}_\alpha$  is the Value-at-Risk at level  $\alpha$ .

**Parameters:**

- `y_true` : array-like, True target values
- `y_pred` : array-like, Predicted values
- `alpha` : float, Percentile threshold (0.90–0.99)

**Returns:**

- `cvar` : float, CVaR metric (lower is better)

**Interpretation:**

- Lower values indicate better tail risk performance
- Industry standard in finance and insurance
- Focuses exclusively on worst-case scenarios

### 4.2.2 Tail Coverage Ratio (TCR)

**Function:** `tailrisk.tail_coverage_ratio(y_true, y_pred, quantile=0.99)`

**Definition:**

TCR measures the fraction of extreme tail value captured by predictions:

$$\text{TCR}_q = \frac{\sum_{i \in \mathcal{T}_q} \hat{y}_i}{\sum_{i \in \mathcal{T}_q} y_i} \quad (6)$$

where  $\mathcal{T}_q$  is the set of indices where  $y_i > q\text{-th percentile}$ .

**Interpretation:**

- $\text{TCR} = 1.0$  : Perfect coverage (100% of tail value captured)
- $\text{TCR} < 1.0$  : Underprediction (dangerous for risk management)
- $\text{TCR} > 1.0$  : Overprediction (conservative, may be acceptable)

#### 4.2.3 Detection Rate

**Function:** `tailrisk.detection_rate(y_true, y_pred, quantile=0.95)`

**Definition:**

Percentage of actual extreme cases correctly predicted as extreme:

$$\text{Detection Rate}_q = \frac{|\{i : y_i > q_y \wedge \hat{y}_i > q_{\hat{y}}\}|}{|\{i : y_i > q_y\}|} \quad (7)$$

where  $q_y$  and  $q_{\hat{y}}$  are the  $q$ -th percentiles of  $y$  and  $\hat{y}$ , respectively.

**Use Case:**

Critical for early warning systems and reserve planning.

#### 4.2.4 tail\_validation\_summary

**Function:** `tailrisk.tail_validation_summary(y_true, y_pred)`

Returns comprehensive dictionary of all tail risk metrics:

- `mse_overall` : Overall mean squared error
- `mse_extreme` : MSE on extreme values ( $> 95$ th percentile)
- `cvar_90`, `cvar_95`, `cvar_99` : CVaR at different thresholds
- `tcr_95`, `tcr_99` : Tail coverage ratios
- `detection_90`, `detection_95`, `detection_99` : Detection rates
- `lar` : Loss-at-Risk metric

## 5 Methodology

### 5.1 The Hybrid Meta-Learning Framework

The HybridMetaLearner implements a novel architecture designed to optimize tail risk prediction:

#### 5.1.1 Stage 1: Diverse Base Models

Train multiple heterogeneous models:

$$\mathcal{M} = \{f_1, f_2, \dots, f_k\} \quad (8)$$

Recommended model types:

- Tree-based: Random Forests, Gradient Boosting
- Linear: Ridge, Lasso
- Ensemble: XGBoost, LightGBM

#### 5.1.2 Stage 2a: Quantile Meta-Model

Use quantile regression on meta-features:

$$\hat{f}_{\text{quantile}}^{(\alpha)} = \arg \min_f \sum_{i=1}^n \rho_\alpha(y_i - f(\mathbf{z}_i)) \quad (9)$$

where  $\rho_\alpha$  is the pinball loss function:

$$\rho_\alpha(u) = u(\alpha - \mathbb{I}(u < 0)) \quad (10)$$

and  $\mathbf{z}_i$  are meta-features from cross-validated base model predictions.

#### 5.1.3 Stage 2b: LaR-Weighted Optimization

Optimize weights for base model combination:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n w_i \left( y_i - \sum_{j=1}^k w_j f_j(\mathbf{x}_i) \right)^2 \quad (11)$$

subject to  $\sum_{j=1}^k w_j = 1$  and  $w_j \geq 0$ .

#### 5.1.4 Stage 3: Hybrid Blending

Final prediction combines both approaches:

$$\hat{y}_{\text{final}} = \lambda \cdot \hat{y}_{\text{quantile}} + (1 - \lambda) \cdot \hat{y}_{\text{LaR}} \quad (12)$$

#### Rationale:

- Quantile model excels at *detecting* extreme events
- LaR model maintains overall *accuracy*
- Blending provides optimal balance (typically  $\lambda = 0.25$ )

## 6 Use Cases

### 6.1 Insurance Claims Prediction

#### Problem Statement:

Catastrophic claims (top 1% of cases) often represent 30–50% of total losses. Traditional models severely underestimate these claims, leading to:

- Inadequate reserves
- Insolvency risk
- Incorrect premium pricing

#### TailRisk Solution:

```

1 from tailrisk import HybridMetaLearner
2 from sklearn.ensemble import RandomForestRegressor
3
4 # Insurance-specific configuration
5 model = HybridMetaLearner(
6     base_estimators=[
7         ('rf', RandomForestRegressor(n_estimators=200)),
8         # Add domain-specific models
9     ],
10    quantile=0.99,      # Focus on top 1% of claims
11    blend_lambda=0.20   # Prioritize accuracy
12)
13
14 model.fit(X_train, y_train)
15 claim_predictions = model.predict(X_test)

```

#### Results:

Typical improvements over baseline models:

- TCR@99% improves from 0.058 to 0.102 (+76%)
- CVaR(95%) reduces by 10–15%
- Detection Rate@95% increases from 0.8% to 2.5%

### 6.2 Financial Risk Management

#### Problem Statement:

Value-at-Risk (VaR) models systematically underestimate tail risk, as evidenced during the 2008 financial crisis. Regulatory frameworks (Basel III) now require CVaR-based risk assessment.

#### TailRisk Solution:

```

1 from tailrisk import LaRRegressor
2 from tailrisk.metrics import cvar_loss
3
4 # Financial loss prediction
5 model = LaRRegressor(alpha=2.5) # Higher focus on tails
6 model.fit(historical_features, historical_losses)
7
8 # Predict potential losses
9 portfolio_risk = model.predict(current_positions)
10
11 # Evaluate CVaR for regulatory compliance
12 cvar_95 = cvar_loss(actual_losses, predictions, alpha=0.95)

```

**Impact:**

- Improved capital allocation
- Regulatory compliance (Basel III)
- Better stress testing capabilities

### 6.3 Healthcare Cost Prediction

**Problem Statement:**

Rare catastrophic cases (ICU admissions, complex surgeries) drive healthcare costs but are poorly predicted by standard models, leading to:

- Inaccurate premium pricing
- Unsustainable risk pools
- Unexpected losses for insurers

**TailRisk Solution:**

Focus on identifying and accurately pricing high-cost cases:

```
1 from tailrisk import HybridMetaLearner
2 from tailrisk.utils import print_tail_validation
3
4 model = HybridMetaLearner(
5     base_estimators=medical_models,
6     quantile=0.95,
7     blend_lambda=0.25
8 )
9
10 model.fit(patient_features, medical_costs)
11 cost_predictions = model.predict(new_patients)
12
13 # Evaluate tail performance
14 print_tail_validation(
15     actual_costs,
16     cost_predictions,
17     model_name="Healthcare Costs"
18 )
```

## 7 Model Comparison Example

### 7.1 Complete Workflow

```
1 import numpy as np
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LinearRegression
4 from sklearn.ensemble import (
5     RandomForestRegressor,
6     GradientBoostingRegressor
7 )
8 from tailrisk import LaRRegressor, HybridMetaLearner
9 from tailrisk.utils import compare_models, print_tail_validation
10
11 # Generate heavy-tailed data
12 np.random.seed(42)
13 n_samples = 5000
14 X = np.random.randn(n_samples, 10)
15
16 # Exponential distribution (heavy tail)
17 y = np.random.exponential(scale=1000, size=n_samples)
18
19 # Add catastrophic claims
20 n_catastrophic = int(0.05 * n_samples)
21 catastrophic_idx = np.random.choice(
22     n_samples,
23     n_catastrophic,
24     replace=False
25 )
26 y[catastrophic_idx] = np.random.exponential(
27     scale=20000,
28     size=n_catastrophic
29 )
30
31 # Split data
32 X_train, X_test, y_train, y_test = train_test_split(
33     X, y, test_size=0.2, random_state=42
34 )
35
36 # Train baseline model
37 baseline = LinearRegression()
38 baseline.fit(X_train, y_train)
39 y_pred_baseline = baseline.predict(X_test)
40
41 # Train LaR model
42 lar_model = LaRRegressor(alpha=2.0)
43 lar_model.fit(X_train, y_train)
44 y_pred_lar = lar_model.predict(X_test)
45
46 # Train Hybrid Meta-Learner
47 base_models = [
48     ('rf', RandomForestRegressor(
49         n_estimators=100,
50         max_depth=10,
51         random_state=42
52     )),
53     ('gb', GradientBoostingRegressor(
54         n_estimators=100,
55         max_depth=5,
56         random_state=42
57 ))]
```

```

57     ))
58 ]
59
60 hybrid_model = HybridMetaLearner(
61     base_estimators=base_models,
62     quantile=0.95,
63     blend_lambda=0.25,
64     lar_alpha=1.5,
65     cv_folds=5
66 )
67
68 hybrid_model.fit(X_train, y_train)
69 y_pred_hybrid = hybrid_model.predict(X_test)
70
71 # Compare all models
72 results = compare_models(y_test, {
73     'Baseline (LR)': y_pred_baseline,
74     'LaR (alpha=2.0)': y_pred_lar,
75     'Hybrid Meta': y_pred_hybrid
76 })
77
78 # Detailed validation for best model
79 print_tail_validation(
80     y_test,
81     y_pred_hybrid,
82     model_name="Hybrid Meta-Learner"
83 )

```

## 7.2 Expected Output

Table 1: Model Comparison Results

Metric	Baseline	LaR	Hybrid
MSE (Overall)	2,808,591	2,815,649	2,979,450
MSE (Extreme)	260,918,453	259,101,161	252,549,889
CVaR (95%)	3,146	3,086	2,824
Detection @ 90%	3.1%	3.2%	6.1%
Detection @ 95%	0.8%	0.9%	2.5%
TCR @ 95%	0.165	0.166	0.306
TCR @ 99%	0.058	0.059	0.102

### Key Observations:

1. Hybrid model shows modest increase in overall MSE (+6%)
2. **Extreme MSE improves by 3%** (critical for tail risk)
3. **CVaR improves by 10%** (better worst-case performance)
4. **TCR@99% nearly doubles** (+76%, from 0.058 to 0.102)
5. **Detection rates triple** (0.8% → 2.5% @ 95%)

### Interpretation:

The Hybrid Meta-Learner trades a small increase in average error for dramatic improvements in tail risk metrics. This is the desired behavior for applications where extreme events drive losses.

## 8 Best Practices

### 8.1 Model Selection

Table 2: Model Selection Guide

Model	When to Use	Advantages
LaRRegressor	<ul style="list-style-type: none"> <li>• Simple problems</li> <li>• Limited data</li> <li>• Need interpretability</li> </ul>	Fast training, interpretable, good baseline
HybridMetaLearner	<ul style="list-style-type: none"> <li>• Complex patterns</li> <li>• Sufficient data</li> <li>• Critical applications</li> </ul>	Best tail performance, robust, production-ready
CVaRWeightedEnsemble	<ul style="list-style-type: none"> <li>• Multiple models</li> <li>• Model combination</li> <li>• Ensemble methods</li> </ul>	Automatic weighting, leverages diversity

### 8.2 Hyperparameter Tuning

#### 8.2.1 LaRRegressor Alpha Parameter

- $\alpha \in [1.0, 1.5]$  : Moderate tail focus
- $\alpha \in [1.5, 2.5]$  : Strong tail focus (recommended)
- $\alpha > 2.5$  : Very aggressive (may sacrifice average performance)

#### 8.2.2 HybridMetaLearner Configuration

- **quantile:**
  - 0.90 for moderate tail focus
  - 0.95 for standard tail risk (recommended)
  - 0.99 for extreme tail events only
- **blend\_lambda:**
  - 0.15–0.25: Prioritize overall accuracy
  - 0.25–0.30: Balanced (recommended)
  - 0.30–0.40: Prioritize tail detection
- **Base Estimators:**
  - Minimum 2, recommended 3–5
  - Include diverse types (tree, linear, boosting)
  - Avoid too many correlated models

## 8.3 Data Considerations

### 8.3.1 Minimum Sample Size

Table 3: Recommended Sample Sizes

Model	Minimum Samples
LaRRegressor	200–500
HybridMetaLearner	1,000+
CVaRWeightedEnsemble	500–1,000

### 8.3.2 Train/Test Split Strategy

For tail risk applications, use **stratified splitting** to ensure test set contains representative extreme values:

```

1 from sklearn.model_selection import StratifiedKFold
2
3 # Create binary indicator for tail events
4 threshold = np.percentile(y, 95)
5 is_tail = (y >= threshold).astype(int)
6
7 # Stratified split
8 splitter = StratifiedKFold(
9     n_splits=5,
10    shuffle=True,
11    random_state=42
12)
13
14 for train_idx, test_idx in splitter.split(X, is_tail):
15     X_train, X_test = X[train_idx], X[test_idx]
16     y_train, y_test = y[train_idx], y[test_idx]
17     break # Use first fold

```

## 8.4 Evaluation

Always evaluate using tail-specific metrics:

1. Primary metrics: CVaR, TCR
2. Secondary metrics: Detection Rate, LaR
3. For context: Overall MSE, MAE

Never rely on MSE/MAE alone for tail risk applications!

## 9 Scikit-learn Integration

TailRisk is fully compatible with scikit-learn ecosystem:

### 9.1 GridSearchCV

```

1 from sklearn.model_selection import GridSearchCV
2 from sklearn.metrics import make_scorer
3 from tailrisk import LaRRegressor, cvar_loss
4
5 # Custom scorer for CVaR
6 def cvar_scorer(y_true, y_pred):
7     return -cvar_loss(y_true, y_pred, alpha=0.95)
8
9 cvar_score = make_scorer(cvar_scorer, greater_is_better=True)
10
11 # Grid search
12 param_grid = {
13     'alpha': [1.0, 1.5, 2.0, 2.5, 3.0]
14 }
15
16 grid = GridSearchCV(
17     LaRRegressor(),
18     param_grid,
19     cv=5,
20     scoring=cvar_score,
21     verbose=1
22 )
23
24 grid.fit(X_train, y_train)
25 best_model = grid.best_estimator_

```

### 9.2 Pipeline

```

1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from tailrisk import LaRRegressor
4
5 # Create pipeline
6 pipe = Pipeline([
7     ('scaler', StandardScaler()),
8     ('model', LaRRegressor(alpha=2.0))
9 ])
10
11 # Train
12 pipe.fit(X_train, y_train)
13
14 # Predict
15 predictions = pipe.predict(X_test)

```

### 9.3 Cross-Validation

```

1 from sklearn.model_selection import cross_val_score
2
3 scores = cross_val_score(
4     LaRRegressor(alpha=2.0),

```

```
5     X, y,
6     cv=5,
7     scoring='neg_mean_squared_error',
8 )
9
10 print(f"CV MSE: {-scores.mean():.2f} +/- {scores.std():.2f}")
```

## 10 Advanced Topics

### 10.1 Custom Base Estimators

Create domain-specific base models:

```

1 from sklearn.base import BaseEstimator, RegressorMixin
2
3 class DomainSpecificModel(BaseEstimator, RegressorMixin):
4     """Custom model incorporating domain knowledge"""
5
6     def __init__(self, domain_param=1.0):
7         self.domain_param = domain_param
8
9     def fit(self, X, y, sample_weight=None):
10        # Implement custom fitting logic
11        # Must support sample_weight for LaR
12        return self
13
14    def predict(self, X):
15        # Implement prediction logic
16        return predictions
17
18 # Use in HybridMetaLearner
19 base_models = [
20     ('rf', RandomForestRegressor()),
21     ('custom', DomainSpecificModel(domain_param=2.0))
22 ]
23
24 model = HybridMetaLearner(base_estimators=base_models)

```

### 10.2 Ensemble Stacking

Combine TailRisk models with traditional stacking:

```

1 from sklearn.ensemble import StackingRegressor
2 from tailrisk import LaRRegressor, HybridMetaLearner
3
4 # Create stacking ensemble
5 stacking = StackingRegressor(
6     estimators=[
7         ('lar', LaRRegressor(alpha=2.0)),
8         ('hybrid', HybridMetaLearner(
9             base_estimators=base_models
10            )),
11        ],
12        final_estimator=Ridge(alpha=1.0)
13    )
14
15 stacking.fit(X_train, y_train)

```

### 10.3 Production Deployment

#### 10.3.1 Model Serialization

```

1 import joblib
2
3 # Save model

```

```
4 joblib.dump(hybrid_model, 'production_model.pkl')
5
6 # Load model
7 loaded_model = joblib.load('production_model.pkl')
8
9 # Use in production
10 predictions = loaded_model.predict(new_data)
```

### 10.3.2 Monitoring

Monitor model performance in production:

```
1 def monitor_tail_performance(
2     y_true,
3     y_pred,
4     alert_threshold_tcr=0.15
5 ):
6     """Monitor production model performance"""
7     from tailrisk import tail_coverage_ratio
8
9     tcr = tail_coverage_ratio(
10         y_true,
11         y_pred,
12         quantile=0.99
13     )
14
15     if tcr < alert_threshold_tcr:
16         print(f"WARNING: TCR dropped to {tcr:.3f}")
17         print("Model may need retraining!")
18         # Trigger alert/notification
19
20     return tcr
21
22 # Use monthly or quarterly
23 tcr_current = monitor_tail_performance(
24     recent_actuals,
25     recent_predictions
26 )
```

## 11 Troubleshooting

### 11.1 Common Issues

#### 11.1.1 Issue: Negative Predictions

**Problem:** Model predicts negative values for inherently positive targets (costs, claims).

**Solution:**

```

1 # Post-process predictions
2 predictions = np.maximum(predictions, 0)
3
4 # Or use log-transformation
5 from sklearn.compose import TransformedTargetRegressor
6
7 model = TransformedTargetRegressor(
8     regressor=LARRegressor(alpha=2.0),
9     func=np.log1p,
10    inverse_func=np.expm1
11)

```

#### 11.1.2 Issue: Poor Tail Performance

Possible causes and solutions:

##### 1. Insufficient tail samples

- Solution: Collect more data, especially extreme cases
- Consider synthetic oversampling of tail events

##### 2. Alpha too low

- Solution: Increase  $\alpha$  to 2.5 or 3.0

##### 3. Wrong base models

- Solution: Include non-linear models (trees, boosting)
- Ensure model diversity

#### 11.1.3 Issue: Training Too Slow

**Solutions:**

```

1 # Reduce cv_folds
2 model = HybridMetaLearner(
3     base_estimators=base_models,
4     cv_folds=3 # Instead of 5
5 )
6
7 # Use fewer/simpler base models
8 base_models = [
9     ('rf', RandomForestRegressor(n_estimators=50)), # Reduce trees
10    ('ridge', Ridge()) # Fast linear model
11 ]
12
13 # Parallelize base model training
14 model = RandomForestRegressor(
15     n_estimators=100,

```

```
16     n_jobs=-1 # Use all CPU cores  
17 )
```

## 11.2 Debugging

Enable detailed output:

```
1 import logging  
2  
3 logging.basicConfig(level=logging.DEBUG)  
4  
5 # Training will show detailed progress  
6 model.fit(X_train, y_train)
```

## 12 Theoretical Background

### 12.1 Quantile Regression

Quantile regression estimates conditional quantiles of the response variable:

$$Q_\tau(Y|X = x) = \arg \min_q \mathbb{E}[\rho_\tau(Y - q)|X = x] \quad (13)$$

where the check function (pinball loss) is:

$$\rho_\tau(u) = u(\tau - \mathbb{I}(u < 0)) = \begin{cases} \tau \cdot u & \text{if } u \geq 0 \\ (\tau - 1) \cdot u & \text{if } u < 0 \end{cases} \quad (14)$$

For tail risk, we focus on high quantiles ( $\tau = 0.95, 0.99$ ) to capture extreme outcomes.

### 12.2 Extreme Value Theory

Heavy-tailed distributions (common in insurance and finance) satisfy:

$$\lim_{x \rightarrow \infty} \frac{\mathbb{P}(X > x + y | X > x)}{\mathbb{P}(X > y)} = 1 \quad (15)$$

This memoryless property makes standard regression challenging. TailRisk's weighted approaches account for this heavy-tail behavior.

### 12.3 CVaR as Coherent Risk Measure

CVaR satisfies the four axioms of coherent risk measures:

1. **Monotonicity:** If  $X \leq Y$ , then  $\text{CVaR}(X) \leq \text{CVaR}(Y)$
2. **Translation invariance:**  $\text{CVaR}(X + c) = \text{CVaR}(X) + c$
3. **Homogeneity:**  $\text{CVaR}(\lambda X) = \lambda \cdot \text{CVaR}(X)$  for  $\lambda \geq 0$
4. **Subadditivity:**  $\text{CVaR}(X + Y) \leq \text{CVaR}(X) + \text{CVaR}(Y)$

This makes CVaR superior to VaR for risk management applications.

## 13 References

### 13.1 Package Documentation

- GitHub Repository: <https://github.com/VishalLakshmiNarayanan/TailriskLib>
- PyPI Page: <https://pypi.org/project/tailrisk/>
- API Reference: [https://github.com/VishalLakshmiNarayanan/TailriskLib/blob/main/docs/API\\_REFERENCE.md](https://github.com/VishalLakshmiNarayanan/TailriskLib/blob/main/docs/API_REFERENCE.md)
- Tutorial: <https://github.com/VishalLakshmiNarayanan/TailriskLib/blob/main/docs/TUTORIAL.md>

### 13.2 Related Software

- Scikit-learn: <https://scikit-learn.org/>
- StatsModels: <https://www.statsmodels.org/> (Quantile Regression)
- PyRisk: <https://github.com/quantopian/pyfolio> (Financial Risk Analysis)

## 14 Citation

If you use TailRisk in academic work, please cite:

```
@software{tailrisk2025,
  author = {Vishal Lakshmi Narayanan},
  title = {TailRisk: Risk-Aware Machine Learning
            for Tail Risk Modeling},
  year = {2025},
  version = {0.1.2},
  url = {https://github.com/VishalLakshmiNarayanan/TailriskLib},
  doi = {10.5281/zenodo.XXXXXXX}
}
```

## 15 License

TailRisk is released under the MIT License:

MIT License

Copyright (c) 2025 Vishal Lakshmi Narayanan

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

## 16 Contact & Support

### 16.1 Author

#### Vishal Lakshmi Narayanan

- Email: [lvishal1607@gmail.com](mailto:lvishal1607@gmail.com)
- GitHub: [@VishalLakshmiNarayanan](https://github.com/VishalLakshmiNarayanan)
- Repository: <https://github.com/VishalLakshmiNarayanan/TailriskLib>

### 16.2 Bug Reports

Report bugs at: <https://github.com/VishalLakshmiNarayanan/TailriskLib/issues>

Please include:

- Python version
- TailRisk version
- Minimal reproducible example
- Expected vs. actual behavior
- Full error traceback

### 16.3 Feature Requests

Submit feature requests as GitHub issues with:

- Use case description
- Proposed API (if applicable)
- Relevance to tail risk modeling

### 16.4 Contributing

Contributions are welcome! See `CONTRIBUTING.md` in the repository for:

- Development setup
- Code style guidelines
- Pull request process
- Testing requirements